

Divvy-Biking Beyond Boundaries

November 25, 2023

Seasons of Cycling: Analyzing Divvy's Year-Long Data Trends

1 Data Collection and Preparation

Divvy Trip Data The datasets were downloaded from this [link](#). A total of 12 csv files (1 file per month) were upload into Python as Pandas Dataframes. The files were combined into 1 file using `.concat()` method.

```
[1]: #import packages
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

#read each csv file
nov_22 = pd.read_csv("202211-divvy-tripdata.csv")
dec_22 = pd.read_csv("202212-divvy-tripdata.csv")
jan_23 = pd.read_csv("202301-divvy-tripdata.csv")
feb_23 = pd.read_csv("202302-divvy-tripdata.csv")
mar_23 = pd.read_csv("202303-divvy-tripdata.csv")
apr_23 = pd.read_csv("202304-divvy-tripdata.csv")
may_23 = pd.read_csv("202305-divvy-tripdata.csv")
jun_23 = pd.read_csv("202306-divvy-tripdata.csv")
jul_23 = pd.read_csv("202307-divvy-tripdata.csv")
aug_23 = pd.read_csv("202308-divvy-tripdata.csv")
sep_23 = pd.read_csv("202309-divvy-tripdata.csv")
oct_23 = pd.read_csv("202310-divvy-tripdata.csv")
print('import done')

#use concat to combine 12 csv
df=pd.
    ↪concat([nov_22,dec_22,jan_23,feb_23,mar_23,apr_23,may_23,jun_23,jul_23,aug_23,sep_23,oct_23,
    ↪ignore_index=True)

#drop unnecessary columns
#df.drop(df.columns[[5, 7]], axis=1, inplace = True)

#inspect dataframe
```

```
#df.head()
df.info()

#df.describe()
```

```
import done
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5652827 entries, 0 to 5652826
Data columns (total 13 columns):
#   Column                Dtype
---  -
0   ride_id               object
1   rideable_type         object
2   started_at           object
3   ended_at             object
4   start_station_name    object
5   start_station_id     object
6   end_station_name     object
7   end_station_id       object
8   start_lat            float64
9   start_lng            float64
10  end_lat              float64
11  end_lng              float64
12  member_casual        object
dtypes: float64(4), object(9)
memory usage: 560.7+ MB
```

The combined dataframe has 5 million records (5,652,827) and has 13 columns (attributes). Upon closer inspection it is observed that the `started_at` and `ended_at` columns are of incorrect datatype so we have converted them back to datetime datatype using pandas `to_datetime()` method.

```
[2]: df['started_at'] = pd.to_datetime(df['started_at'])
     df['ended_at'] = pd.to_datetime(df['ended_at'])
```

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5652827 entries, 0 to 5652826
Data columns (total 13 columns):
#   Column                Dtype
---  -
0   ride_id               object
1   rideable_type         object
2   started_at           datetime64[ns]
3   ended_at             datetime64[ns]
4   start_station_name    object
5   start_station_id     object
6   end_station_name     object
7   end_station_id       object
```

```

8   start_lat      float64
9   start_lng      float64
10  end_lat        float64
11  end_lng        float64
12  member_casual  object
dtypes: datetime64[ns](2), float64(4), object(7)
memory usage: 560.7+ MB

```

2 Data Exploration and Cleaning

Each record in the dataframe has `ride_id` along with time, station information like name, location(latitude and longitude), bike type and user_type for each ride recorded. **ride_id**, Let's talk about `ride_id` for a sec. Think of it as the VIP pass for each record in our data party – it's unique for every guest and there are no plus-ones. The cool thing? Every `ride_id` is like a secret code, exactly 16 characters long. We've checked them all, and they're in perfect shape. So, guess what? We don't need to fuss over this column anymore. It's all good to go as is!

```

[4]: # Calculate % unique values per column
duplicates = df.nunique().reset_index()
duplicates.columns = ['column', 'unique_values']
duplicates['unique%'] = round((duplicates['unique_values'] / len(df)) * 100, 2)

# Calculate % missing values per column
missing = df.isna().sum().reset_index()
missing.columns = ['column', 'missing_values']
missing['missing%'] = round((missing['missing_values'] / len(df)) * 100, 2)

# Combine the dataframes
combined_df = pd.merge(duplicates, missing, on='column')
print(combined_df)

```

	column	unique_values	unique%	missing_values	missing%
0	ride_id	5652827	100.00	0	0.00
1	rideable_type	3	0.00	0	0.00
2	started_at	4764981	84.29	0	0.00
3	ended_at	4776456	84.50	0	0.00
4	start_station_name	1579	0.03	866243	15.32
5	start_station_id	1494	0.03	866375	15.33
6	end_station_name	1589	0.03	918796	16.25
7	end_station_id	1503	0.03	918937	16.26
8	start_lat	784998	13.89	0	0.00
9	start_lng	745056	13.18	0	0.00
10	end_lat	13873	0.25	6759	0.12
11	end_lng	13990	0.25	6759	0.12
12	member_casual	2	0.00	0	0.00

```
[5]: #checking if all the ride_id have exactly 16 charaters
ride_id_len = (df['ride_id'].str.len()==16).all()
print(ride_id_len)

#count duplicates on unique column
# this is done to find if there are any duplicate records present in the data as
↳ride_id is the primary key, it is chosen
print('Total duplicates in ride_id column: ',df['ride_id'].duplicated().sum())
```

True

Total duplicates in ride_id column: 0

rideable_type, this column is like the ID badge for each bike, telling us what model was used for the ride. We've got three types listed: classic_bike, docked_bike, and electric_bike. But here's an interesting update: we realized that 'docked_bike' is actually an old name for what we now call a 'classic_bike.' So, we decided to give our data a little makeover. We've updated 'docked_bike' to 'classic_bike' across the board, and this tweak has brought a new lease of life to 86098 records. It's all about keeping things consistent and clear!

```
[6]: #finding the unique values in rideable_type column
print(df['rideable_type'].unique())
```

['electric_bike' 'classic_bike' 'docked_bike']

```
[7]: #counting the number of records that contained docked_bike
count_before_change = (df['rideable_type'] == 'docked_bike').sum()
print('The number of records having docked_bike type before:',
↳count_before_change)
#changing docked_bike to classic_bike
df['rideable_type'] = df['rideable_type'].replace('docked_bike','classic_bike')
#counting the number of records after change
count_after_change = (df['rideable_type'] == 'docked_bike').sum()
print('The number of records having docked_bike type after the change:',
↳count_after_change)
#checking for null values or empty values in rideable_type
print('Total number of empty values or null values in rideable_type :
↳',df['rideable_type'].isna().sum())
```

The number of records having docked_bike type before: 86098

The number of records having docked_bike type after the change: 0

Total number of empty values or null values in rideable_type : 0

started_at and **ended_at**, these columns are our timekeepers in the dataset. They don't just tell us when each bike trip kicked off and wrapped up; they're the key to unlocking much more. With these timestamps, we can calculate the duration of each ride – a vital piece of the puzzle. But there's more: by breaking down these dates and times, we can see patterns based on the day of the week and specific dates. This slicing and dicing of time not only makes our data richer for analysis but also adds a dash of life to our data visualizations. The granularity we get from these details is super valuable, helping us spot trends and derive insights that would otherwise be hidden in broader data.

```
[8]: #we have already converted the started_at and ended_at columns to datetime
      ↪format and hence we can proceed to find the duration of the ride
df['date'] = df['started_at'].dt.date
df['day'] = df['started_at'].dt.day_name()
df['ride_duration'] = ((df['ended_at'] - df['started_at']).dt.total_seconds() /
      ↪60)
df.head()
```

```
[8]:
```

	ride_id	rideable_type	started_at	ended_at	\
0	BCC66FC6FAB27CC7	electric_bike	2022-11-10 06:21:55	2022-11-10 06:31:27	
1	772AB67E902C180F	classic_bike	2022-11-04 07:31:55	2022-11-04 07:46:25	
2	585EAD07FDEC0152	classic_bike	2022-11-21 17:20:29	2022-11-21 17:34:36	
3	91C4E7ED3C262FF9	classic_bike	2022-11-25 17:29:34	2022-11-25 17:45:15	
4	709206A3104CABC8	classic_bike	2022-11-29 17:24:25	2022-11-29 17:42:51	

	start_station_name	start_station_id	end_station_name	\
0	Canal St & Adams St	13011	St. Clair St & Erie St	
1	Canal St & Adams St	13011	St. Clair St & Erie St	
2	Indiana Ave & Roosevelt Rd	SL-005	St. Clair St & Erie St	
3	Indiana Ave & Roosevelt Rd	SL-005	St. Clair St & Erie St	
4	Indiana Ave & Roosevelt Rd	SL-005	St. Clair St & Erie St	

	end_station_id	start_lat	start_lng	end_lat	end_lng	member_casual	\
0	13016	41.879401	-87.639848	41.894345	-87.622798	member	
1	13016	41.879255	-87.639904	41.894345	-87.622798	member	
2	13016	41.867888	-87.623041	41.894345	-87.622798	member	
3	13016	41.867888	-87.623041	41.894345	-87.622798	member	
4	13016	41.867888	-87.623041	41.894345	-87.622798	member	

	date	day	ride_duration
0	2022-11-10	Thursday	9.533333
1	2022-11-04	Friday	14.500000
2	2022-11-21	Monday	14.116667
3	2022-11-25	Friday	15.683333
4	2022-11-29	Tuesday	18.433333

2.0.1 Dealing with outliers

Here's a curious thing we spotted: among the sea of rides, some are as brief as under a minute, while others stretch beyond 24 hours – talk about extremes! We've decided to label these ultra-short and ultra-long rides as outliers. It's like finding a needle in a haystack, but we did it – and to keep our data neat and tidy, we're going to remove these outliers. A total of 156,036 rows, to be exact, are saying goodbye to our main dataset. But don't worry, they're not going into the data void; we're giving them a new home in a separate dataframe, `df_duration_noise`. This move helps us focus on the more typical rides and maintain the integrity of our analysis.

```
[9]: # Count rows before filtering
count_before = len(df)

# First, create a DataFrame of outliers
df_duration_noise = df[(df['ride_duration'] < 1) | (df['ride_duration'] > 24*60)]

# Then, filter df to remove outliers
df = df[(df['ride_duration'] >= 1) & (df['ride_duration'] <= 24*60)]

# Count rows after filtering
count_after = len(df)

# Print the number of rows deleted
print('This change has resulted in deleting', count_before - count_after,
      'rows')

# Check the shape of df_duration_noise
df_duration_noise.shape
```

This change has resulted in deleting 156036 rows

```
[9]: (156036, 16)
```

start_station_name and end_station_name We've noticed a bit of a puzzle: quite a few trips are missing either their starting or ending station names. Now, here's where things get interesting. For classic bikes, it's a must to have both a start and an end at a docking station. But electric bikes? They're the free spirits of our dataset – they can end their journeys pretty much anywhere, no dock required. So, to keep our data tidy and meaningful, we've made a decision: any classic bike records missing station names are going to be moved to a new home, a separate dataframe we're calling `df_station_noise`. This way, we keep our main dataset clean and focused on the complete journeys.

```
[10]: # Identify classic bike records missing station names
classic_missing_stations = (df['rideable_type'] == 'classic') &
    (df['start_station_name'].isna() | df['end_station_name'].isna())

# Create df_station_noise DataFrame
df_station_noise = df[classic_missing_stations].copy()

# Update the main DataFrame by removing these records
df = df[~classic_missing_stations]
```

```
[11]: #storing the outliers or noise in a dataframe called df_noise, this is done to
    preserve the data integrity and future analysis
df_noise = pd.concat([df_duration_noise, df_station_noise])
df_noise.head()
```

```
[11]:
```

	ride_id	rideable_type	started_at	ended_at	\
149	7F8CA9B17D7E2B5F	classic_bike	2022-11-03 11:52:03	2022-11-03 11:52:47	
151	9320FCC9994902BC	electric_bike	2022-11-08 05:17:18	2022-11-08 05:17:21	
152	E7372C2C8A9BFCA7	electric_bike	2022-11-08 05:16:41	2022-11-08 05:16:45	
189	2B096F11BFFAEEF4	electric_bike	2022-11-25 10:47:39	2022-11-25 10:48:05	
412	61A73ABE32A0FFE6	classic_bike	2022-11-03 15:49:17	2022-11-03 15:49:19	

	start_station_name	start_station_id	end_station_name	\
149	Desplaines St & Kinzie St	TA1306000003	Desplaines St & Kinzie St	
151	Hoyne Ave & Balmoral Ave	655	Hoyne Ave & Balmoral Ave	
152	Hoyne Ave & Balmoral Ave	655	Hoyne Ave & Balmoral Ave	
189	NaN	NaN	Ashland Ave & Lake St	
412	Walsh Park	18067	Walsh Park	

	end_station_id	start_lat	start_lng	end_lat	end_lng	member_casual	\
149	TA1306000003	41.888716	-87.644448	41.888716	-87.644448	member	
151	655	41.979913	-87.682015	41.979851	-87.681932	member	
152	655	41.979833	-87.682010	41.979851	-87.681932	member	
189	13073	41.890000	-87.670000	41.885920	-87.667170	member	
412	18067	41.914610	-87.667968	41.914610	-87.667968	member	

	date	day	ride_duration
149	2022-11-03	Thursday	0.733333
151	2022-11-08	Tuesday	0.050000
152	2022-11-08	Tuesday	0.066667
189	2022-11-25	Friday	0.433333
412	2022-11-03	Thursday	0.033333

```
[12]: df_noise.shape
```

```
[12]: (156036, 16)
```

```
[13]: #checking member_casual column for possible values
print((df['member_casual']).unique())
df['member_casual'].isna().sum()
```

```
['member' 'casual']
```

```
[13]: 0
```

In our journey through the data, we've come across a neat little detail about the `member_casual` column. It turns out, it's pretty straightforward – just two types of riders here, member and casual. And guess what? There's not a single null or empty spot in sight for this column.

But here's where it gets a bit more complex: the `start_station_name`, `start_station_id`, `end_station_name`, and `end_station_id` columns are a different story. They've got a fair share of nulls and empties. However, we've decided not to show these records the exit door. Why? Because these gaps actually tell us something important – they reflect the unique flexibility of electric bikes, which don't always need a specific docking station. So, instead of dropping this valuable info,

we've taken a creative turn: we're labeling these unknowns with an unknown value. This way, we acknowledge the gaps without losing the bigger picture of our bike-riding saga.

```
[14]: print("finding the number of missing values and their percentage after removing
↳the outliers and noise")
missing = df.isna().sum().reset_index()
missing.columns = ['column', 'missing_values']
missing['missing%'] = round((missing['missing_values'] / len(df)) * 100, 2)
print(missing)
```

finding the number of missing values and their percentage after removing the outliers and noise

	column	missing_values	missing%
0	ride_id	0	0.00
1	rideable_type	0	0.00
2	started_at	0	0.00
3	ended_at	0	0.00
4	start_station_name	821489	14.94
5	start_station_id	821614	14.95
6	end_station_name	853288	15.52
7	end_station_id	853426	15.53
8	start_lat	0	0.00
9	start_lng	0	0.00
10	end_lat	801	0.01
11	end_lng	801	0.01
12	member_casual	0	0.00
13	date	0	0.00
14	day	0	0.00
15	ride_duration	0	0.00

```
[15]: # Columns to replace missing values with 'unknown'
columns_to_replace = ['start_station_name', 'start_station_id',
↳'end_station_name', 'end_station_id']

# Replace NaN values in each specified column with 'unknown'
for column in columns_to_replace:
    df[column] = df[column].fillna('unknown')
```

3 Data Analysis

Now moving on to the data visualization using the processed data to derive insights.

```
[16]: ride_counts = df['member_casual'].value_counts()

# Data for the pie chart
labels = ride_counts.index
sizes = ride_counts.values
```




```

# Plotting the pie chart
plt.figure(figsize=(4, 3))
plt.pie(sizes, labels=labels, autopct='%1.1f%%', startangle=140)
plt.axis('equal') # Ensures the pie chart is circular
plt.title('Distribution of Rides: Members vs. Casual Riders')
plt.show()

```

Distribution of Rides: Members vs. Casual Riders



```

[17]: average Ride Duration = df.groupby('member_casual')['ride_duration'].mean()
print(average Ride Duration)
average Ride Duration = {'Member': 21.198452, 'Casual': 12.371359} # in minutes
'''

# Data preparation
categories = list(average Ride Duration.keys())
values = list(average Ride Duration.values())

# Creating the bar chart
plt.figure(figsize=(8, 6))
plt.bar(categories, values, color=['blue', 'green'])
plt.xlabel('Rider Type')
plt.ylabel('Average Ride Duration (minutes)')
plt.title('Average Ride Duration: Member vs Casual Riders')
plt.xticks(categories)
plt.show()
'''

```

```

member_casual
casual    21.198452
member    12.371359
Name: ride_duration, dtype: float64

```

```
[17]: "\n\n# Data preparation\nncategories = list(average Ride Duration.keys())\nvalues
= list(average Ride Duration.values())\n\n# Creating the bar
chart\nplt.figure(figsize=(8, 6))\nplt.bar(categories, values, color=['blue',
'green'])\nplt.xlabel('Rider Type')\nplt.ylabel('Average Ride Duration
(minutes)')\nplt.title('Average Ride Duration: Member vs Casual
Riders')\nplt.xticks(categories)\nplt.show()\n"
```

```
[18]: # Define a function to categorize days into 'Weekday' and 'Weekend'
def categorize_day(day):
    if day.weekday() < 5: # 0 to 4 corresponds to Monday to Friday
        return 'Weekday'
    else: # 5 and 6 corresponds to Saturday and Sunday
        return 'Weekend'


# Apply the function to create a new column
df['day_type'] = df['date'].apply(categorize_day)

# Group by rider type and day type, then calculate the mean duration
average_duration = df.groupby(['member_casual', 'day_type'])['ride_duration'].
    .mean().reset_index()

# Pivot the data for easier plotting
pivot_data = average_duration.pivot(index='day_type', columns='member_casual',
    values='ride_duration')
```

```
[19]: # Plotting the data
pivot_data.plot(kind='bar', figsize=(10, 6))

plt.xlabel('Day Type')
plt.ylabel('Average Ride Duration')
plt.title('Average Ride Duration by Rider Type and Day Type')
plt.xticks(rotation=0) # Rotate x-axis labels to show them horizontally
plt.show()
```



```
[20]: df['date'] = pd.to_datetime(df['date'])
# Create a new column for the day of the week
df['day_of_week'] = df['date'].dt.day_name()

# Group by day of the week and rider type, then count the trips
trips_by_day_rider = df.groupby(['day_of_week', 'member_casual']).size().
    ↪reset_index(name='trip_count')


# Pivot the data for easier plotting
pivot_data = trips_by_day_rider.pivot(index='day_of_week',
    ↪columns='member_casual', values='trip_count')

# Ensure the days are ordered
days_order = ["Monday", "Tuesday", "Wednesday", "Thursday", "Friday",
    ↪"Saturday", "Sunday"]
pivot_data = pivot_data.reindex(days_order)

# Plotting the data
pivot_data.plot(kind='line', marker='o', figsize=(10, 6))

plt.xlabel('Day of the Week')
plt.ylabel('Number of Trips')
plt.title('Number of Trips by Day of the Week and Rider Type')
```

```
plt.grid(True)
plt.xticks(range(len(days_order)), days_order) # Set x-ticks to days of the week
plt.show()
```



From the graph, we can infer the following:

Casual Riders:

The number of trips by casual riders is higher than that of members on every day of the week. Casual ridership appears to peak midweek, with the highest number on Wednesday, and then gradually declines towards the weekend.

Member Riders:

The pattern for member riders is quite different. The number of trips starts low on Monday, increases significantly on Tuesday, remains relatively steady through Friday, and then spikes on Saturday. The number of trips for members drops on Sunday, indicating perhaps a lesser preference for using the service on that day compared to Saturday.

Overall Trends:

Casual riders seem to use the service more consistently across the week with a peak in the middle of the week. Members show a preference for using the service towards the end of the workweek and on Saturdays.

These observations can suggest several underlying behaviors and preferences:

Casual riders might include tourists or occasional users who are more active during the week, possibly indicating leisure or errand-related activities that are not tied to the workweek schedule.


Member riders may use the service for commuting purposes, reflected in the higher number of trips on weekdays, especially towards the latter half of the week and on Saturdays for weekend activities.

The drop in member trips on Sunday might indicate a day of rest or non-reliance on bike-sharing services, possibly due to reduced work or social activities.


```
[21]: '''  
      !pip install pivottablejs  
      from pivottablejs import pivot_ui  
      pivot_ui(df.head(20000))  
      pivot_ui(df_noise)'''
```

```
[21]: '\n!pip install pivottablejs\nfrom pivottablejs import  
pivot_ui\npivot_ui(df.head(20000))\npivot_ui(df_noise)'
```

```
[22]: import seaborn as sns  
  
      # Creating the heatmap  
      plt.figure(figsize=(10, 6))  
      sns.heatmap(pivot_data, annot=True, fmt="d", cmap='YlGnBu', linewidths=.5)  
  
      plt.title('Heatmap of Number of Trips by Day of the Week and Rider Type')  
      plt.ylabel('Day of the Week')  
      plt.xlabel('Rider Type')  
      plt.show()
```



The heatmap illustrates the distribution of Divvy bike-sharing trips across different days of the week for casual and member riders. Members show a consistent increase in trips as the week progresses, peaking on Tuesday, while casual riders' trips peak on Saturdays. The heatmap's color gradient indicates that members generally take more trips than casual riders on weekdays, with the highest volume on Tuesday.

```
[23]: import pandas as pd
      from scipy import stats

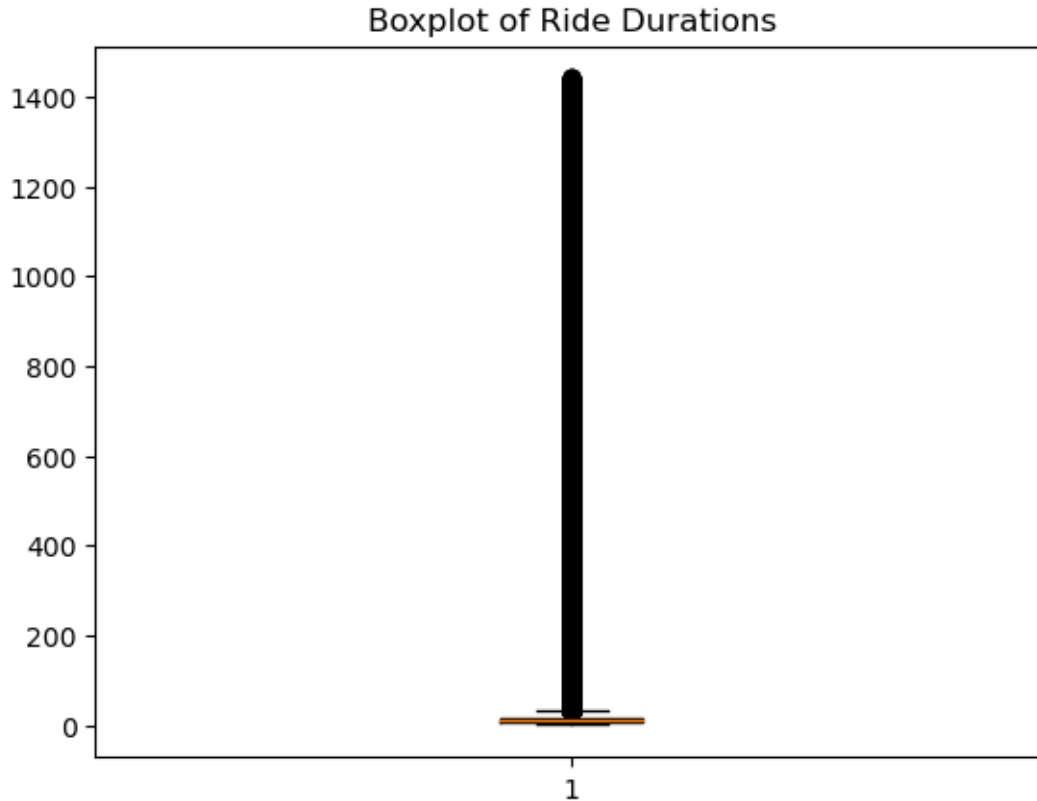
      # Assuming df is your DataFrame
      z_scores = stats.zscore(df['ride_duration'])
      outliers_z = df[(z_scores < -3) | (z_scores > 3)]
```

```
[24]: Q1 = df['ride_duration'].quantile(0.25)
      Q3 = df['ride_duration'].quantile(0.75)
      IQR = Q3 - Q1

      outliers_iqr = df[(df['ride_duration'] < (Q1 - 1.5 * IQR)) |
      ↪ (df['ride_duration'] > (Q3 + 1.5 * IQR))]
```

```
[25]: import matplotlib.pyplot as plt

      plt.boxplot(df['ride_duration'])
      plt.title('Boxplot of Ride Durations')
      plt.show()
```



z-score outlier detection and IQR outlier detection are statistical methods used to identify abnormal points in the dataset, specifically within the ride_duration variable:

Z-Score Outlier Detection:

The z-score method standardizes the entire dataset by converting data points into z-scores, which represent the number of standard deviations a point is from the mean. This project uses a threshold of 3 standard deviations to identify outliers, meaning any ride duration more than 3 standard deviations from the mean ride duration is considered an outlier. This method is sensitive to the mean and standard deviation, and therefore can be affected by extreme values or a non-normal distribution of data.

IQR Outlier Detection:

The IQR method involves calculating the interquartile range, which is the range between the first quartile (25th percentile) and the third quartile (75th percentile) of the data. Outliers are defined as observations that fall below $Q1 - 1.5IQR$ or above $Q3 + 1.5IQR$. This does not assume a normal distribution of the data and is less influenced by extreme values. In this project, any ride_duration outside of these bounds is classified as an outlier and is a candidate for exclusion from further analysis to prevent skewing the results.

Boxplot Visualization:

The boxplot created in the project provides a visual representation of the distribution of


ride_duration. It displays the median, quartiles, and potential outliers, which are individual points that appear outside the whiskers of the boxplot (typically set at $1.5 \times \text{IQR}$ from the quartiles). This visualization aids in confirming the presence of outliers and understanding the spread and symmetry of the data.

Both z-scores and IQR are used here to rigorously identify ride durations that are unusually long or short, which might otherwise bias the analysis. The boxplot serves as a visual confirmation of these findings, offering a clear picture of the data distribution and highlighting any potential outliers.

```
[26]: #time series analysis
#This plots the average ride duration for each month to observe any trends over
      ↪time.
import matplotlib.pyplot as plt
import seaborn as sns
#df['month'] = df['started_at'].dt.to_period('M')
# Ensure 'ride_duration' is numeric
df['ride_duration'] = pd.to_numeric(df['ride_duration'], errors='coerce')

# Convert 'started_at' to Period (monthly), then to string for plotting
df['month'] = df['started_at'].dt.to_period('M').astype(str)


monthly_duration = df.groupby('month')['ride_duration'].mean().reset_index()
plt.figure(figsize=(12, 6))
sns.lineplot(data=monthly_duration, x='month', y='ride_duration')
plt.title('Average Ride Duration Per Month')
plt.xlabel('Month')
plt.ylabel('Average Ride Duration (minutes)')
plt.xticks(rotation=45)
plt.show()
```



The line graph depicts the average ride duration per month for Divvy bike-sharing from November 2022 to October 2023. There is a noticeable dip in ride durations in December 2022, after which there's a steady increase, peaking in July 2023. Following this peak, there's a sharp decline in August and September, with a slight recovery in October. This trend may suggest seasonal patterns in ride usage, with longer rides in the warmer months and shorter rides in the colder months.

The observed pattern likely reflects user behavior changes in response to weather conditions, with shorter rides during the cold winter months and longer rides during the warm summer months, indicating a preference for using bike-sharing services for longer periods when the weather is more favorable. The decline in late summer could be due to a return to school or work routines, suggesting a shift in the reasons for bike usage. The data could be vital for Divvy in planning resource allocation, and maintenance schedules to match these seasonal trends.

```
[27]: #rider behaviour analysis
#This will compare the ride durations between members and casual riders across_
↳ days of the week.
df['day_of_week'] = df['started_at'].dt.day_name()
plt.figure(figsize=(12, 6))
sns.barplot(data=df, x='day_of_week', y='ride_duration', hue='member_casual')
plt.title('Ride Duration Comparison by Rider Type Across Days of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Average Ride Duration (minutes)')
plt.show()
```



```
[28]: #Geospatial analysis
import folium
from folium.plugins import HeatMap

# Create a base map
map = folium.Map(location=[df['start_lat'].mean(), df['start_lng'].mean()],
    ↪zoom_start=12)

# Add a heatmap to the base map
HeatMap(data=df[['start_lat', 'start_lng']].dropna(), radius=10).add_to(map)

map
```

```
[28]: <folium.folium.Map at 0x20994e24be0>
```